

3. Fitting a linear model to a sample by least squares

- Recall the sample version of the linear model. Data are y_1, y_2, \dots, y_n and on each individual i we have p explanatory variables $x_{i1}, x_{i2}, \dots, x_{ip}$.

$$(LM1) \quad y_i = b_1 x_{i1} + b_2 x_{i2} + \dots + b_p x_{ip} + e_i \quad \text{for } i = 1, 2, \dots, n$$

- Using summation notation, we can equivalently write

$$(LM2) \quad y_i = \sum_{j=1}^p x_{ij} b_j + e_i \quad \text{for } i = 1, 2, \dots, n$$

- We can also use matrix notation. Define column vectors $\mathbf{y} = (y_1, y_2, \dots, y_n)$, $\mathbf{e} = (e_1, e_2, \dots, e_n)$ and $\mathbf{b} = (b_1, b_2, \dots, b_p)$. Define the matrix of explanatory variables, $\mathbb{X} = [x_{ij}]_{n \times p}$. In matrix notation, writing (LM1) or (LM2) is exactly the same as

$$(LM3) \quad \mathbf{y} = \mathbb{X} \mathbf{b} + \mathbf{e}$$

- Matrices give a compact way to write the linear model, and also a good way to carry out the necessary computations.

The least squares formula

- We seek the **least squares** choice of \mathbf{b} that minimizes the sum of squared error, $\sum_{i=1}^n e_i^2$.
- Since n is usually much bigger than p , there is usually no value of \mathbf{b} for which we can exactly explain the data using the explanatory matrix \mathbb{X} .
- In other words, there is no choice of \mathbf{b} which solves $\mathbb{X}\mathbf{b} = \mathbf{y}$.
- The least squares choice of \mathbf{b} turns out to be

$$(LM4) \quad \mathbf{b} = (\mathbb{X}^T \mathbb{X})^{-1} \mathbb{X}^T \mathbf{y}$$

- We will check that this is the formula R uses to fit a linear model.
- We will also gain understanding of (LM4) by studying the **simple linear regression** model $y_i = b_1 x_i + b_2 + e_i$ for which $p = 2$.
- In the simple linear regression model, b_1 and b_2 are called the slope and the intercept. In general, b_1, \dots, b_p are called the **coefficients** of the linear model. We call \mathbf{b} the coefficient vector.
- In R, we obtain \mathbf{b} using the `coef()` function as demonstrated below.

Checking the coefficient estimates from R

- Consider the example from Chapter 1, where `L_detrended` is life expectancy for each year, after subtracting a linear trend, and `U_detrended` is the corresponding detrended unemployment.

```
lm1 <- lm(L_detrended~U_detrended)
coef(lm1)
```

```
## (Intercept) U_detrended
##      0.2899928      0.1313673
```

- Now, we can construct the X matrix corresponding to this linear model and ask R to compute the coefficients using the formula (LM4).

```
X <- cbind(U_detrended, intercept=rep(1,length(U_detrended)))
```

```
solve( t(X) %*% X ) %*% t(X) %*% L_detrended
```

```
##                [,1]
## U_detrended 0.1313673
## intercept   0.2899928
```

Checking the \mathbf{X} matrix we constructed

- The matrix calculation on the previous slide matches the coefficients produced by `lm()`.
- Take some time to check that our R implementation matches the formula (LM4).
- We're fairly sure we got the computation right, because we exactly matched `lm()`, but it is a good idea to look at the \mathbf{X} matrix we constructed.

```
head(X)
```

```
##      U_detrended intercept
## 1  -1.0075234         1
## 2   1.1027941         1
## 3   0.4881116         1
## 4  -1.5349043         1
## 5  -1.8662535         1
## 6  -2.0059360         1
```

```
length(U_detrended)
```

```
## [1] 68
```

```
dim(X)
```

```
## [1] 68  2
```

Naming the \mathbb{X} matrix in the linear model $\mathbf{y} = \mathbb{X}\mathbf{b} + \mathbf{e}$

- “The \mathbb{X} matrix” is not a great name since we would have the same model if we had called it \mathbb{Z} .
- Many names are used for \mathbb{X} for the many different purposes of linear models.
- Sheather’s textbook calls \mathbb{X} the **matrix of predictor variables** or **matrix of explanatory variables**.
- We call \mathbb{X} the **design matrix** in situations where x_{ij} is the setting of adjustable variable j for the i th run of an experiment. For example, y_i could be the strength of an alloy made up of a fraction x_{ij} of metal j for $j = 1, \dots, p - 1$. We would also want to include an intercept, $x_{ip} = 1$.
- \mathbb{X} can also be called the **matrix of covariates**.
- Sometimes, \mathbf{y} is called the **dependent variable** and \mathbb{X} is the **matrix of independent variables**. Scientifically, an independent variable is one that can be set by the scientist. However, independence has a different technical meaning in statistics.

Fitted values

- The **fitted values** are the estimates of the data based on the explanatory variables. For our linear model, these fitted values are

$$\hat{y} = b_1x_{i1} + b_2x_{i2} + \cdots + b_px_{ip}, \quad \text{for } i = 1, 2, \dots, n.$$

- The vector of least squares fitted values $\hat{\mathbf{y}} = (\hat{y}_1, \dots, \hat{y}_n)$ is given by

$$(LM5) \quad \hat{\mathbf{y}} = \mathbb{X}\mathbf{b} = \mathbb{X}(\mathbb{X}^T\mathbb{X})^{-1}\mathbb{X}^T\mathbf{y}.$$

- It is worth checking we now understand how R produces the fitted values for predicting detrended life expectancy using unemployment:

```
my_fitted_values<-X %*% solve(t(X)%*%X) %*% t(X) %*% L_detrended
```

```
lm1$fitted.values[1:2]
```

```
## [1] 0.1576371 0.4348639
```

```
my_fitted_values[1:2]
```


```
## [1] 0.1576371 0.4348639
```

- We see that the matrix calculation (LM5) exactly matches the fitted values of the `lm1` model that we built earlier using `lm()`.

Plotting the data

- We have already seen plots of the life expectancy and unemployment data before. When you fit a linear model you should look at the data and the fitted values. We plot the fitted values two different ways.

```
plot(L_detrended~U_detrended)
lines(U_detrended,my_fitted_values,lty="solid",col="blue")
abline(coef(lm1),lty="dotted",col="red",lwd=2)
```



Question 3.1. Learn about the `abline()` and `lines()` functions. Explain to yourself why the solid blue line and the dotted red line coincide.

Deriving the formula for the least squares coefficient vector

This material will not be tested in the exam. It is presented to show you an application of differentiation and to explain where the formula (LM4) for \mathbf{b} comes from.

- We derive (LM4) for the simple linear regression model (SLR1).
- The **sum of squared error** is also the **sum of the squared residuals** and is known as the **residual sum of squares (RSS)**. For simple linear regression, this is

$$\text{RSS} = \sum_{i=1}^n (y_i - ax_i - b)^2$$

- To find a and b minimizing RSS, we differentiate with respect to a and b and set the derivatives equal to zero.

Question. Check that

$$\frac{\partial}{\partial a} \text{RSS} = 2a \sum_{i=1}^n x_i^2 - 2 \sum_{i=1}^n x_i (y_i - b)$$